

# Getting started with Docker

In this tutorial you'll learn:

- How to run **containers**.
- How to define and build **images**.
- How to create and use **volumes**.
- How to define and use **networks**.

## Prerequisites:

- Having installed Docker on your computer or having imported the Linux virtual machine either with VirtualBox or with Multipass.

## Related course material:

- Slides of the first lecture. You can find them on Edunao.
- Handbook (poly) of the first lecture. You can find it on Edunao.

Using a virtual machine with Multipass? [Click here for more info](#)

### Multipass commands

- To start the virtual machine, type `multipass start cloudvm`.
- The folder `/home/ubuntu/labs` in the virtual machine is mounted (i.e., linked) to the folder on YOUR computer (let's call it **LAB**) where you'll store all your lab material. You specified this folder when you installed the virtual machine.
- If you don't remember the path to the folder linked to `/home/ubuntu/labs`, type the command `multipass info cloudvm`.
- To open a terminal into the virtual machine, type `multipass shell cloudvm`.
- When the terminal opens, the current working directory is `/home/ubuntu`.
- Type `cd labs` to move to the folder where you'll find your lab material.
- You'll use the virtual machine terminal only to type the Docker commands. You can use all the tools installed on your machine to create and manage the files in the directory **LAB** directly from your computer.
- Once the lab is over, type `exit` to leave the virtual machine terminal. This will lead you back to the terminal of your computer.
- In the terminal of your computer, type `multipass stop cloudvm` to stop the virtual machine. This will NOT destroy your files! It just stops the virtual machine from needlessly using the resources of your computer.

### Terminology

- You'll use the **terminal** to run Docker commands. The terminal is the **client** that communicates with the Docker daemon.
- Docker runs **containers** on your computer. We'll refer to your computer as the **host**, the containers being the **guests**.
- A **containerized application** is an application running in a container.

# 1 Running containers

We now learn how to use the command `docker run` and some of its options. In the following exercises, we'll run containers from the image named *alpine* that is available on the DockerHub registry. This image provides a lightweight distribution (i.e., it doesn't contain many features) of Linux.

## Exercise

**Exercise 1.1.** You want to run the container from the latest version of the image *alpine*. Which command would you write in the terminal? Execute it.

Solution

The goal of this exercise is to start playing with the `docker run` command. Since the question doesn't say anything about the options, nor does it mention the command to run inside the container, we'd type:

```
docker run alpine
```

## Exercise

**Exercise 1.2.** Is the container still running?

Solution

In order to list all containers still running on the host, type the following command:

```
docker container ls
```

Your container shouldn't appear in the output, because it's not running. In order to see all containers, including those that are not running, type the following command:

```
docker container ls -a
```

## Exercise

**Exercise 1.3.** By looking at the command executed within the container (`/bin/sh`), can you tell why the container stopped without giving any output?

Solution

The command is `/bin/sh`; the container runs a Linux terminal. But since we didn't specify what to do with that terminal (we didn't run any Linux command, nor we tried to access the terminal), the container stopped.

We're now going to do something useful with the image *alpine*. Make sure you read the **good practices** that you should adopt while playing with images and containers.

### Good practices

1. **Name your containers.** Although Docker assigns a default name to a new container, it's usually a good practice to give a container a name of your choice to make it easily distinguishable. You can do it by using the option `--name`. Try the following:

```
docker run --name my-alpine alpine
```

As before, the container stops immediately. If you list all your containers by typing again:

```
docker container ls -a
```

you should see a container named *my-alpine*.

2. **Remove automatically a container if you use it once.** Unless you want to reuse your container later, you can ask Docker to automatically remove it when it stops by using the option `--rm`. This will prevent unused containers from taking up too much disk space.

Try the following:

```
docker run --rm --name container-to-remove alpine
```

If you list all the containers you should see that there is no container named *container-to-remove*.

3. **Remove unused containers.** Stopped containers that have been run without using the option `--rm` are still stored in the host. If you want to remove a specific container (e.g., *my-alpine*), use the following command:

```
docker container rm my-alpine
```

If you want to remove all stopped containers, use the following command:

```
docker container prune
```

4. **Remove unused images.** Images can take up a lot of disk space. As a result, you should remember to remove those that you don't intend to use any longer. The commands to remove a specific image and prune unused ones are `docker image rm` and `docker image prune -a` respectively.

## 1.1 Pass a command to the containerized application

Remember that the template of `docker run` is the following:

```
docker run [options] image-name [command] [arg]
```

The optional parameter *command* refers to a command that you can pass the containerized application, possibly with some arguments (parameter *arg*).

Let's see an example. As we saw before, when we run a container from the image *alpine*, a Linux terminal `/bin/sh` is launched.

### Notice

The Linux terminal `/bin/sh` is run within the container. Henceforth, we'll use the following terms:

- **Host terminal.** The terminal that you use to interact with the operating system of your computer.
- **Guest terminal.** The terminal that is run within the container.

By using the optional parameter *command*, we can run a command in the guest terminal.

### Exercise

**Exercise 1.4.** Run a container from the image *alpine* and execute the Linux command `ls` that lists the content of the current directory.

- Where are the listed files stored? In the host or in the container?

Solution

```
docker run --rm --name ls-test alpine ls
```

- The command `ls` is run in the guest terminal, therefore what we see in the output is a list of files stored in the container.

### Notice

In Exercise 1.4 the command `ls` is executed in the guest terminal, but its output is redirected to the host terminal.

In other words, when we run the container, we don't interact directly with the guest terminal; we just send a command and the output is redirected to the host terminal.

Now let's see how to execute a command in the guest terminal that also requires an argument.

### Exercise

**Exercise 1.5.** By using the Linux utility `ping`, check whether the Web site `www.centralesupelec.fr` is reachable.

Solution

```
docker run --rm --name ping-test alpine ping www.centralesupelec.fr
```

In order to interrupt `ping` just type the key combination that you use to interrupt any other command in your terminal. (typically `Ctrl-C` on Windows and `Cmd-C` in MacOS).

Now run the following command:

```
docker run --name my-alpine -it alpine
```

**Note:** we didn't use the option `--rm` (the container will not be removed when we stop it, we're going to use it again). Moreover, we didn't specify any command to run in the guest terminal.

### Exercise

**Exercise 1.6.** What do you obtain?

Solution

When we run a container from the image `alpine`, the command `/bin/sh` is executed within the container. Since we specified the option `-it`, what we obtain is an access to the Linux terminal running in the container.

## 1.2 Starting and stopping containers.

`docker run` is a shorthand for two Docker commands, namely `docker create`, that creates a container from an image, and `docker start`, that starts the container after its creation.

Suppose now that you want to download a Web page by using Linux Alpine. You can use the Linux command `wget` followed by the URL of the page that you want to download.

### Exercise

**Exercise 1.7.** By using the guest terminal in the container *my-alpine*, download this Web page.

- Where will the Web page be saved? The host computer or the container?

Solution

Just type in *my-alpine* guest terminal the following command:

```
wget https://www.centralesupelec.fr/fr/presentation
```

- The Web page will be saved in the current directory of the container. You can verify that the file is there by typing `ls` in the guest terminal.

Want to copy the downloaded file from the container to your computer?

To copy the file `presentation` to your working directory, type the following command in the host terminal:

```
docker cp <containerId>:/presentation .
```

In the previous command, replace `<containerId>` with the identifier of your container. You can obtain the identifier of the container from the output of the command `docker container ls -a`.

### Where are the containers and image files stored?

If you use MacOS or Windows

The files managed by Docker are not stored directly in your computer, but in the Linux virtual machine installed and operated by Docker Desktop (remember, Docker always need Linux to be executed).

Therefore, you need to open a terminal inside that Linux virtual machine by typing the following command:

```
docker run -it --rm --privileged --pid=host justincormack/nsenter1
```

Once the terminal is opened, you can follow the instructions given below for Linux users.

If you use Docker on Linux

- All files managed by Docker are stored under folder `/var/lib/docker`.
- To access that folder, you need to be root (i.e., administrator). Type the command `sudo su`.
- If you type `ls /var/lib/docker` you can look at the folders stored under this directory. You'll see that there are folders corresponding to the different objects managed by Docker (containers, images, volumes and networks).
- To locate the files of a specific container, you first need to get the **container identifier** by typing `docker container ls -a`.
- Type the command `docker inspect <container-id>` (replace `<container-id>` with the identifier of the container that you intend to inspect).
- Locate the field `UpperDir`. The value of this field is the path to the directory (let's call it `CONTAINER_DIR`) that contains the upper layer of the container (the writable layer). It should be a path that ends by `/diff`.
- If you type `cd CONTAINER_DIR` (replace `CONTAINER_DIR` with the value of the field `UpperDir`) you can finally see the files stored in your container.

In *my-alpine* guest terminal type `exit`. This closes the guest terminal and, as a result, stops the container.

### NOTICE

Stopping the container will not erase any of the files stored in the container. Removing the container will.

If you want to start the container *my-alpine* again, you can use the following command:

```
docker container start -ai my-alpine
```

This will open the guest terminal of the container again; type `ls` to verify that the Web page that you downloaded before is still there.

Homework (optional)

Suppose that you need to download all the figures of this Web page. The Linux utility `wget` comes in handy. However, you don't have Linux and you'd like to avoid the hassle of installing it on your computer, or in a virtual machine, just for this task.

A great alternative is to run Linux in a Docker container. Unfortunately, the Alpine distribution that we've been playing with doesn't provide an implementation of `wget` with all the options that we need.

We turn to another Linux distribution, **Ubuntu**, for which DockerHub has several images.

### Exercise

**Exercise 1.8.** Run a container with Ubuntu (latest) and open a guest terminal. Call the container *dl-figures*, and avoid the option `--rm`, we'll use this container later.

Solution

```
docker run --name dl-figures -it ubuntu
```

From now on, we'll be interacting with the guest Ubuntu terminal. If you type the command `wget`, you'll get an error (`bash: wget: command not found`).

### Notice

The image *Ubuntu* doesn't include all the commands that you'd find in a full-blown Ubuntu distribution; the reason is to keep the size of the image small, a necessary constraint given that images are transferred over the Internet.

Luckily, there's a way to install `wget` in our Ubuntu distribution. Ubuntu provides a powerful command-line package manager called **Advanced Package Tool** (APT). First, you need to run the following command:

```
apt update
```

which fetches the available packages from a list of sources available in file `/etc/apt/sources.list`.

Then, you can install `wget` by running the following command:

```
apt install -y wget
```

In order to obtain all the figures from a Web page, type the following command:

```
wget -nd -H -p -P /my-figures -A jpg,jpeg,png,gif -e robots=off -w 0.5 https://www.centralesupelec.fr/f
```

You should see in the current directory a new folder named *my-figures* containing the downloaded figures; verify it by typing `ls my-figures`.

Before terminating, don't forget to read your fortune cookie. In the shell, run the following command:

```
apt-get install -y fortune
```

and then:

```
/usr/games/fortune -s
```

When you're done, you can simply type the command `exit` to quit the guest terminal and stop the container.

## 2 Creating Images

A Docker image can be thought of as a template to create and run a container. An image is a file that contains a **layered filesystem** with each layer being **immutable**; this means that the files that belong to a layer cannot be modified or deleted, nor can files be added to a layer.

When a container is created from an image, it will be composed of all the image read-only layers and, on top of them, a writable layer (termed the **container layer**), where all the new files created in the container will be written. For example, the Web page that you downloaded in Exercise 1.7 were stored in the writable layer of that container.

### 2.1 Dockerfiles

The most common way to create an image is to use a **Dockerfile**, a text file that contains all the instructions necessary to build the image. The advantage of the Dockerfile is that it can be interpreted by the Docker engine, which makes the creation of images an automated and repeatable task.

Suppose that we want to create a containerized application to download figures from a Web page. As a template for this application, we need to build a new image, that we'll call *fig-downloader*.

The Dockerfile containing the instructions to build the image *fig-downloader* is as follows:

```
FROM ubuntu
RUN apt-get update
RUN apt-get install -y wget
RUN mkdir -p /my-figures
WORKDIR /my-figures
ENTRYPOINT ["wget", "-nd", "-r", "-A", "jpg,jpeg,bmp,png,gif"]
CMD ["https://www.centralesupelec.fr/fr/presentation"]
```

Here's the explanation:

1. We use the image *ubuntu* as the **base image**. This corresponds to the instruction `FROM ubuntu`.
2. We install the utility `wget` in the base image. This corresponds to the instructions `RUN apt-get update` and `RUN apt-get install -y wget`.
3. We create a directory `my-figures` under the root directory of the image. This corresponds to the instruction `RUN mkdir -p /my-figures`.
4. We set the newly created directory `/my-figures` as the **working directory** of the image. This corresponds to the instruction `WORKDIR /my-figures`.
5. We specify the command to be executed when a container is run from this image. This corresponds to the instruction

```
ENTRYPOINT ["wget", "-nd", "-r", "-A", "jpg,jpeg,bmp,png,gif"]
```

This instruction means: execute `wget` with the options `-nd`, `-r`, `-A`; the last option takes a list of file extensions (`jpg,jpeg,bmp,png,gif`) as its argument.

- Remember that the utility `wget` takes the URL of the Web page as an argument. The URL will be specified when we run the container from the image *fig-downloader*. Optionally, we can specify a default argument by using the keyword `CMD`. The meaning of the instruction:

`CMD ["https://www.centralesupelec.fr/fr/presentation"]`

is: if we don't give any URL when we run the container, the figures will be downloaded from `https://www.centralesupelec.fr/fr/presentation`.

### Exercise

**Exercise 2.1.** What's the relation between the Dockerfile lines and the image layers?

Solution

Each line corresponds to a new layer. The first line corresponds to the bottom layer; the last line to the top layer.

### Exercise

**Exercise 2.2.** Could you identify a problem in this Dockerfile? Modify the Dockerfile accordingly.

Solution

When creating an image, we should keep the number of layers relatively small; in fact, the more the layers, the bigger the image will be. Here we create three separate layers with three `RUN` commands; we can simply merge the three layers. The resulting Dockerfile will be:

```
FROM ubuntu
RUN apt-get update && \
    apt-get install -y wget && \
    mkdir -p /my-figures
WORKDIR /my-figures
ENTRYPOINT ["wget", "-nd", "-r", "-A", "jpg,jpeg,bmp,png,gif"]
CMD ["https://www.centralesupelec.fr/fr/presentation"]
```

## 2.2 Building an image

We're now going to build an image from a Dockerfile.

- Create a directory named *fig-downloader* in your computer with a file named *Dockerfile* inside.
- In the *Dockerfile* write the set of instructions that you proposed in Exercise 2.2.
- In the terminal, set the working directory to *fig-downloader*.
- Build an image called *fig-downloader* by executing the following command:

```
docker build -t fig-downloader .
```

The `.` at the end of the command means that the Docker engine will look for a file named *Dockerfile* in the working directory.



### Good to know

If you give the Dockerfile a different name (say, *Dockerfile-fig-downloader*), the command to build the image will be:

```
docker build -t fig-downloader -f Dockerfile-fig-downloader .
```

The option `-f` is used to specify the name of the Dockerfile.

In order to verify that the new image has been created, type the following command:

```
docker images
```

### Exercise

**Exercise 2.3.** Run the following command:

```
docker run --name dl-1 fig-downloader
```

What does it do? Where are the downloaded pictures?

Solution

We downloaded the figures of this page. The downloaded pictures are in the folder */my-figures* of the container *dl-1*.

### Exercise

**Exercise 2.4.** Run the following command:

```
docker run --name dl-2 fig-downloader https://www.centralesupelec.fr/
```

What does it do? Where are the downloaded pictures?

Solution

We downloaded the figures of this page. We basically overwrote the URL specified by the CMD keyword with a new one. The downloaded pictures are in the folder */my-figures* of the container *dl-2*.

## 2.3 Containerized Python application

Download this archive file and unzip it into your working directory. In this archive you'll find:

- A Dockerfile.
- A Python script *main.py* that asks the user to enter the URL and the language of a Web page, and prints the 10 most frequent words occurring in that page.
- A file *requirements.txt* with the list of the Python packages needed to run the given script.

The content of the Dockerfile is as follows:

```
FROM python:3.7-slim
RUN mkdir -p /app
WORKDIR /app
COPY ./main.py ./requirements.txt /app/
RUN pip install -r requirements.txt
ENTRYPOINT ["python", "main.py"]
```

### Exercise

**Exercise 2.5.** Describe what this Dockerfile does.

Solution

- Takes *python:3.7-slim* as the base image.
- Creates a new folder *app* in the image under the root directory.
- Changes the working directory to */app*.
- Copies the files *main.py* and *requirements.txt* from the local computer to the directory */app* in the image.
- Runs the command `pip install` to install the Python libraries specified in the file *requirements.txt*.
- Executes the command `python main.py`.

### Exercise

**Exercise 2.6.** Build an image called *wordfreq* from this Dockerfile.

Solution

```
docker build -t wordfreq .
```

### Exercise

**Exercise 2.7.** Without changing the Dockerfile, rebuild the same image. What do you notice?

Solution

The build is very fast. Since we didn't change the Dockerfile, the image is rebuilt by using the image layers created previously. This is clearly indicated by the word **CACHED** written at each layer. Using the already stored layers is called **build cache**.

### Exercise

**Exercise 2.8.** What happens if you modify a line in the Python script and you rebuild the image?

Solution

Add any instruction at the end of *main.py*, such as:

```
print("Finish!")
```

then rebuild the image. The three bottom layers are not affected by the modification, therefore they benefit from the build cache. Layer 4 is the first affected by the modification. This layer, and those above, need therefore to be rebuilt.

### Exercise

**Exercise 2.9.** Based on the previous considerations, can you tell what's wrong with this Dockerfile? Modify the Dockerfile accordingly and rebuild the image.

Solution

Each time we modify *main.py* and we rebuild the image, the layer 4 and 5 are recreated, meaning that all the Python packages are downloaded and installed. Depending on the size and number of the packages, this might take some while. A better way to structure the Dockerfile is to install the packages before copying the Python script to the image. Here is how we should modify the Dockerfile:

```
FROM python:3.7-slim
RUN mkdir -p /app
WORKDIR /app
COPY ./requirements.txt /app/
RUN pip install -r requirements.txt
COPY ./main.py /app/
ENTRYPOINT ["python", "main.py"]
```

### Exercise

**Exercise 2.10.** Modify *main.py* by adding a new line of code and rebuild the image. What changed?

Solution

The Python packages are not reinstalled, as a result rebuilding the image takes much less time than before.

Play with the application by running the following command:

```
docker run --rm -it wordfreq
```

The containerized application will prompt you to insert the URL of a webpage and the language of the page (in English). The output will be the 20 most used words in the webpage.

## 3 Data Volumes

In Exercise 2.3 you've been asked to run a container named *dl-1* to download some figures from a Web page. The figures were downloaded into the directory */my-figures* of the container. But we left a question unanswered.

**How do we transfer those figures from the container to the host computer?**

The solution is to use **volumes**.

### 3.1 Docker volumes

A volume can be seen as a virtual storage device attached to a container. All files there are written to a volume survive the containerized application that created them. In other words, when a container is destroyed, all the files created by the application in the container remain in the volume.

Let's create a new Docker volume called *data-volume*:

```
docker volume create data-volume
```

#### Good to know (advanced notion)

##### Where the data will be actually stored?

You can inspect the new volume by typing the following command:

```
docker volume inspect data-volume
```

A *mount point* is indicated; that's the folder where the data will be actually stored. If your computer runs Linux, that folder will be available on the host; if your computer runs Windows or MacOS, you'll not find that folder on your computer. Instead, it will be available in the virtual machine that Docker use on MacOS and Windows.

##### Do you want to see the directory? (Instructions for Windows and MacOS)

One way to look into the hidden VM is to run the following containerized application:

```
docker run -it --rm --privileged --pid=host justincormack/nsenter1
```

This application will open a guest terminal into the VM. You can then use the commands `cd` and `ls` to browse to the directory indicated as the mount path of the new volume.

### 3.1.1 Sharing data

A Docker volume can be used to share data between containers.

#### Exercise

**Exercise 3.1.** Run a container from the image *ubuntu*, specifying the options to:

- Remove the container once its execution is over.
- Interact with the guest Linux terminal in the container.
- Mount the volume *data-volume* at the container's directory */data*.

Solution

```
docker run --rm -it -v data-volume:/data ubuntu
```

1. Type the following command in the guest Linux terminal to create a file *test-file.txt* in the directory */data*:

```
echo "This is a new file" > /data/test-file.txt
```

2. Print to the console the content of the file with the following command:

```
cat /data/test-file.txt
```

3. Type `exit` to leave the guest terminal. Since we've specified the option `--rm`, the container is **destroyed**. Now we're going to verify that *test-file.txt* is still accessible.

#### Exercise

**Exercise 3.2.** Run a container from the image *alpine:latest*, specifying the options to:

- Remove the container once its execution is over.
- Interact with the guest Linux terminal in the container.
- Mount the volume *data-volume* to the directory */my-data* of the container.

Solution

```
docker container run --rm -it -v data-volume:/my-data alpine
```

### Exercise

**Exercise 3.3.** Verify that you can read the file *test-file.txt*. Which folder would you look in?

Solution

We need to look in the folder */my-data* because this is where we mounted *data-volume*.  
`cat /my-data/test-file.txt`

Type `exit` to exit the guest terminal and terminate the container.

## 4 Single-Host Networking

In order to let containers communicate and, therefore, co-operate, Docker defines a simple networking model known as the **container network model** (figure 1).

If you use a Linux virtual machine with Multipass

In this section, you'll need to open several terminals in the virtual machine. You can do it easily by using `byobu`, an advanced window manager already available in your virtual machine.

- Just type `byobu` to launch the window manager.
- If you want to open a new terminal, just press F2.
- If you want to switch from a terminal to another, just press F3 (to move to previous terminal) or F4 (to move to next terminal).
- If you want to close a terminal, just type `exit`.
- When you close all terminals, `byobu` will stop executing.




Figure 1: The Docker container network model

### Exercise

**Exercise 4.1.** Describe the output of the following command:

```
docker network ls
```

Solution

The command lists all the networks created by Docker on your computer. For each network, the values of four attributes are shown:

- The identifier.
- The name.
- The driver used by the network.
- The scope of the network (local or global). A local scope means that the network connects containers running on the same host, as opposed to a global scope that means that containers on different hosts can communicate.

Depending on the containers that you used in the past, you might see different networks. However, three networks are worth noting:

- The network named **bridge**, that uses the driver **bridge** and a local scope. By default, any new container is attached to this network.
- The network named **host**, that uses the driver **host** and a local scope. It's used when we want a container to directly use the network interface of the host. It's important to remember that this network should only be used when analyzing the host's network traffic. In the other cases, using this network exposes the container to all sorts of security risks.
- The network named **none**, that uses the driver **null** and a local scope. Attaching a container to this network means that the container isn't connected to any network, and therefore it's completely isolated.

#### Exercise

**Exercise 4.2.** The following command:

```
docker network inspect bridge
```

outputs the configuration of the network **bridge**. By looking at this configuration, can you tell what IP addresses will be given to the containers attached to this network? What's the IP address of the router of this network?

Solution

The information is specified in the field named **IPAM**, more specifically:

- **Subnet** indicates the range of IP addresses used by the network. The value of this field should be 172.17.0.0/16; the addresses range from 172.17.0.1 to 172.17.255.255.
- **Gateway** indicates the IP address of the router of the network. The value should be 172.17.0.1

## 4.1 Creating networks

By default, any new container is attached to the network named *bridge*.

**Exercise 4.3.** Explain why it is not a good practice to attach all our containers to the same network.

Solution

All new containers will be able to communicate over this network. This is not a good idea. If a hacker can compromise any of these containers, s/he might be able to attack the other containers as well. As a rule of thumb, we should attach two containers to the same network **only** on a need-to-communicate basis.

In order to create a new network, you can use the following command:

```
docker network create network_name
```

#### Exercise

**Exercise 4.4.** Create two networks named *buckingham* and *rochefort* that use the driver *bridge* (see figure 1). By using the `docker network inspect` command, look at the IP addresses of the new networks and write them down.

Solution

Just run the following commands:

```
docker network create buckingham
```

```
docker network create rochefort
```

The IP addresses for the network *buckingham* are 172.18.0.0/16 (addresses from 172.18.0.1 to 172.18.255.255); The IP addresses for the network *rochefort* are: 172.19.0.0/16 (assuming that you create *buckingham* before *rochefort*).

The IP addresses may be different on your machines.

#### Exercise

**Exercise 4.5.** Create three containers *athos*, *porthos* and *aramis* and attach them to the two networks *buckingham* and *rochefort* as displayed in figure 1. **The three containers will open a Linux Alpine shell.** You'll need to launch the commands in three separate tabs of your terminal window.

- What will the IP addresses of the three containers be in the two networks? Remember that *porthos* is attached to two networks, therefore it'll have two network interfaces (endpoints) and, as a result, two IP addresses.
- Verify your answers by inspecting the two networks (use the command `docker network inspect`).

Solution

Here are the commands to run *athos* and *aramis* while connecting them to *buckingham* and *rochefort* respectively.

```
docker run --rm -it --name athos --network buckingham alpine
docker run --rm -it --name aramis --network rochefort alpine
```

Here's the command to run *porthos* and attach it to *buckingham*:

```
docker run --rm -it --name porthos --network buckingham alpine
```

The following command attaches *porthos* to the second network *rochefort*:

```
docker network connect rochefort porthos
```

As for the IP addresses, each network has IP addresses in the range 172.x.0.0/16, where x is 18 in the network *buckingham* and 19 in the network *rochefort*. The address 172.x.0.1 is reserved for the router. Therefore, the containers will be assigned IP addresses from 172.x.0.2. In this solution, we created *athos*, *aramis* and *porthos* in this order. Therefore, the IP addresses will be:

- In network *buckingham*:
  - \* *athos*: 172.18.0.2
  - \* *porthos*: 172.18.0.3
- In network *rochefort*:
  - \* *aramis*: 172.19.0.2
  - \* *porthos*: 172.19.0.3

You can actually verify this configuration by inspecting the two networks with the following commands:

```
docker network inspect buckingham
docker network inspect rochefort
```

The IP addresses might be different on your machines.

## 4.2 Communication between containers

Let's see if and when the three containers can communicate.

### Exercise

**Exercise 4.6.** Which containers are able to communicate? Justify your answer.

Solution

The only containers that cannot communicate are *athos* and *aramis*, because they're not connected to the same network.

### Exercise

**Exercise 4.7.** Try to ping *porthos* from *athos* by using its IP address.

- Which IP address of *porthos* would you use?

Solution

We need to use the IP address assigned to the endpoint linking *porthos* to the network *buckingham*, to which *athos* is connected. In our case, this is 172.18.0.3.



**Exercise**

**Exercise 4.8.** Try to ping *porthos* from *athos* by using its name. Do you succeed? Are you surprised?

Solution

We succeed. Indeed, the network *buckingham* provides a DNS server, that can translate names into IP addresses.

You can now exit the three containers.