

# Introduction to Databases

Gianluca Quercini

Laboratoire de Recherche en Informatique  
CentraleSupélec

2019 – 2020



CentraleSupélec

# Distributed Databases

# Distributed Database

- **Single-server database** (Chapter 1):
  - Database on only one machine.
  - All data under the control of one DBMS. 😊
  - Performances of DBMS decrease as data volume increases. 😞
  - Best solution if the scale of data allows it.
- **Distributed database.**
  - Data reside on multiple machines (a.k.a., *nodes*).
  - Each machine is independent of the others (**shared-nothing architecture**).
  - Allows storage and management of large volumes of data. 😊
  - Far more complex than a single-server database. 😞
  - Scale out only if a single-server database is not viable.

# Characteristics of Distributed Databases

- **Data distribution** options: **replication** and **sharding**.
- **Location transparency**: the user queries the data without even knowing that they are distributed.
- **Replication transparency**: consistency of the data that are replicated.
- **Data management functions**: security and concurrent access control.

# Replication




Department		
codeD	nameD	budget
14	Administration	300,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

Department		
codeD	nameD	budget
14	Administration	300,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

Department		
codeD	nameD	budget
14	Administration	300,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

- + **Scalability.** Multiple nodes receive queries on the same tuple.
- + **Latency.** Worldwide database, replica close to the user.
- + **Fault tolerance.** If a node fails, the others can still answer queries.
- **Consistency.** Keep all replicas up-to-date.

# Replication and Consistency




Department		
codeD	nameD	budget
14	Administration	500,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

Department		
codeD	nameD	budget
14	Administration	300,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

Department		
codeD	nameD	budget
14	Administration	300,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

- **Synchronous update.** The update is propagated immediately.
  - + Short **inconsistency window**.
  - Not viable if updates are frequent.
- **Asynchronous update.** The update is propagated at regular intervals.
  - + Best option when updates are frequent.
  - Large inconsistency window.

# Master-Slave Replication




- Write operations: on the master. Writes are propagated to the slaves.
- Read operations: from the master and from the slaves.
- + No write conflicts.
- Single point of failure.
  - If the master is down, write operations are not allowed.
  - Algorithms exist to **elect** a new master.

# Master-Slave Replication – Read conflicts

**Write:** update (Department, budget=500,000)

**Read:** select (Department, budget)



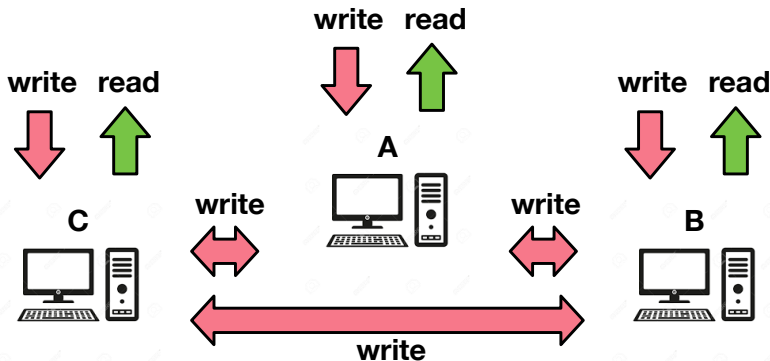
Department		
codeD	nameD	budget
14	Administration	300,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

Department		
codeD	nameD	budget
14	Administration	500,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

Department		
codeD	nameD	budget
14	Administration	500,000
25	Education	150,000
62	Finance	600,000
45	Human Resources	150,000

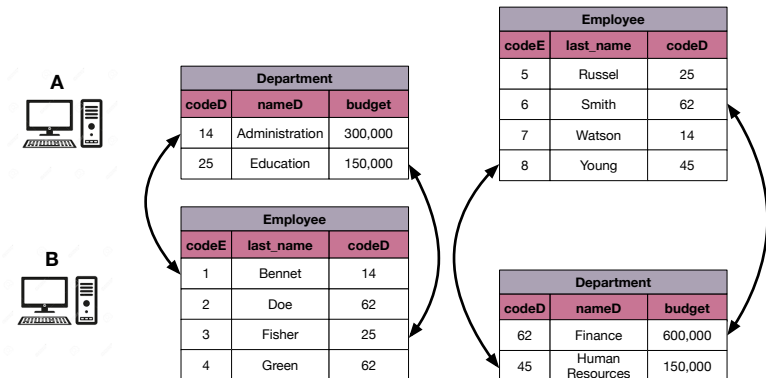


# Peer-to-peer replication



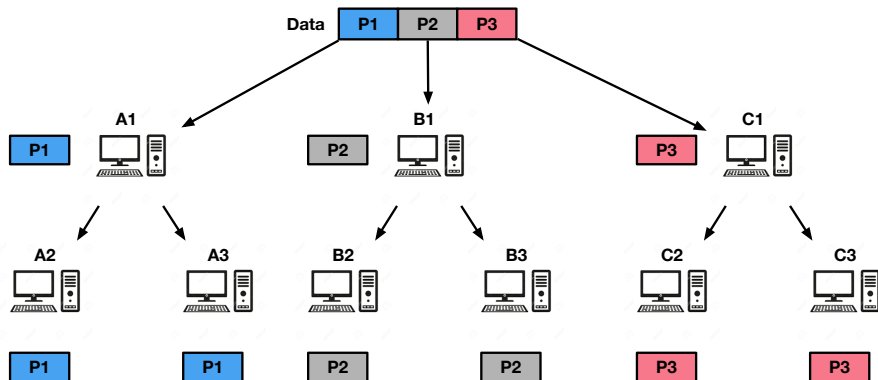
- Write and read operations on any node.
- + No single point of failure (very high availability).
- Write and read conflicts (low consistency).

# Sharding



- Tuples partitioned into balanced **shards**. Shards distributed across the nodes.
  - Loss of data when node fail.
  - Join across different nodes.
  - Updates entail changes in shards.
- + Load balance.
- + No consistency problems.

# Combining Replication and Sharding



# Distributed Transactions

- **Objective:** guarantee ACID properties in distributed databases.
- **Solution:** distributed transactions.
  - sequence of read/write operations that span multiple nodes.
- **Two-phase commit protocol**
  - **Prepare phase.** The coordinator asks all the nodes to prepare to either commit or rollback.
  - **Commit phase.** The coordinator asks the other nodes to commit their operations.
- If only one node fails, the whole transaction is rolled back.
- Distributed transaction management involves a lot of coordination.
  - network traffic
- Ensuring strong consistency in distributed databases is not always a good idea.

# Quorums

- A transactional approach to replication consistency is inefficient. Why?
- **Write quorum.** Number of replicas to lock in a  $n$ -node cluster.


$$QW > \frac{n}{2}$$

- Propagate the update to the remaining  $n - QW$  replicas
- Strong write consistency.
- Weak read consistency.

# Quorums

- **Read quorum.** Number of replicas  $QR$  to read to get a consistent read.

$$QR + QW > n$$



# The CAP Theorem

- **Consistency.** “Equivalent to having a single up-to-date copy of the data” (Brewer).
- **Availability.** A database can perform read/write operations even when some nodes fail.
- **Partition tolerance.** The database can operate when a network partition occurs as if the partition did not happen.

## Theorem (CAP, Brewer 1999)

*Given the three properties of consistency, availability and partition tolerance, a networked shared-data system can have at most two of these properties.*

# The CAP Theorem

Sketch of the proof (Gilbert & Lynch, 2002).

- Suppose that a database is partition tolerant. Two cases when a network partition occurs.
  - Write operations are allowed.
    - Consistency cannot be guaranteed.
    - AP database.
  - Write operations are not allowed.
    - Database unavailable (write operations forbidden on reachable nodes).
    - CP database.
- Consistency and availability only when no network partition occurs.
  - Assuming absence of network partitions means that a database is not partition tolerant.



# The CAP Theorem

- The theorem has been largely misunderstood for years.
- Common interpretation:
  - Network partitions occur.
  - Therefore, the CAP theorem reduces to the choice of consistency over availability.
- However, network partitions are not that frequent.
  - Makes no sense to give up either consistency or availability.
- A distributed database should detect network partitions and operate in a *partition mode* with:
  - Reduced availability (some operations forbidden), or
  - Reduced consistency.
- Partition recovery to resolve the inconsistencies.

# BASE Consistency Model

- ACID transactions are a pessimistic consistency model.
- An optimistic model to consistency, used in distributed databases, is BASE.
- **Basic Availability** (BA). The database appears to work most of the time.
- **Soft state** (S). Write and read inconsistencies can occur.
- **Eventually consistent** (E). The database is consistent at some time.
  - Update propagation.

# NoSQL Databases

# Limitations of the Relational Model

Relational DBMSs have two major limitations:

- **Impedance mismatch.**

- The DBMS always models data as tables.
- An application models data in different ways, depending on their nature.

## Example

- The DBMS models a graph as a collection of tables.
- An application models a graph as an adjacency list.
- Conceived when data distribution was not a concern.
  - Consistency at any cost: not always a good option.
  - Problems when sharding: joins across nodes.
- NoSQL databases conceived to address these concerns.
  - First NoSQL databases: Amazon Dynamo (2007), Google BigTable (2008).

# Characteristics of NoSQL databases

- **NoSQL** means *Not Only SQL* (term coined in 2011).
- NoSQL databases are not based on the relational model.
- NoSQL databases are generally **open-source**.
- NoSQL databases are **cluster-oriented**.
- NoSQL databases tend to privilege **availability** over consistency.
- NoSQL databases are **schemaless**.
- NoSQL databases are classified into four families:
  - **Key-value** databases.
  - **Document** databases.
  - **Column-family** databases.
  - **Graph** databases.
- The first three databases are based on the notion of **aggregate**.

# Aggregate

- **Aggregate:** data structure containing the description of an entity.
- All data in the same aggregate must stay in the same shard.
- Operations within the same aggregate are atomic.
- Schema is flexible and non-normalized.

## Aggregate

```
{
  "codeE": 1,
  "first": "Joseph",
  "last": "Bennett",
  "position": "Office assistant",
  "salary": 55,000,
  "department": [
    {
      "codeD": 14,
      "nameD": "Administration",
      "budget": 30000
    }
  ]
}

{
  "codeE": 2,
  "first": "John",
  "last": "Doe",
  "salary": 45,000
  "department": [
    {
      "codeD": 14,
      "nameD": "Administration",
      "budget": 30000
    }
  ]
}
```

# Aggregate

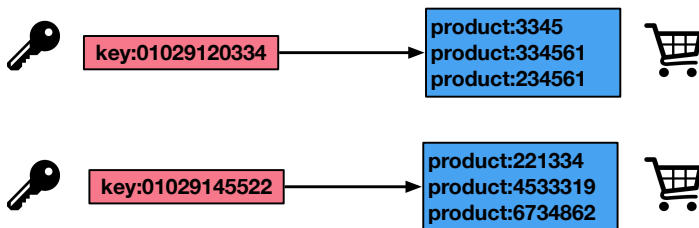
- Different ways to model the data.

## Aggregate

```
{
  "codeD": 14,
  "name": "Administration",
  "budget": 30000,
  "employees": [
    {
      "codeE": 1,
      "first": "Joseph",
      ....
    },
    {
      "codeE": 7,
      "first": "Michael",
      ....
    }
  ]
}
```

# Key-value Data Model

- Data are modeled as **key-value pairs**.
  - **Key**: alphanumeric string auto-generated by the database.
  - **Value**: an aggregate.
  - **Query**: get a value given its key.
- Fast read/write operations.
- Little to no checks on integrity constraints.
- Example: shopping cart.
- Key-value databases: Amazon Dynamo, Voldemort, Riak, Memcached DB.



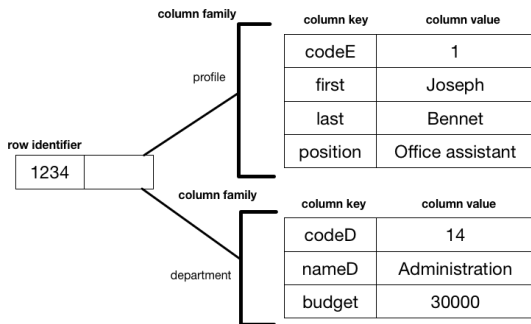


# Document Data Model

- Data are modeled as **key-value pairs**.
  - **Key**: alphanumeric string auto-generated by the database.
  - **Value**: an aggregate (called **document**).
  - **Query**: get documents by key and by the values of their properties.
- Example in detail: MongoDB (Chapter 4).

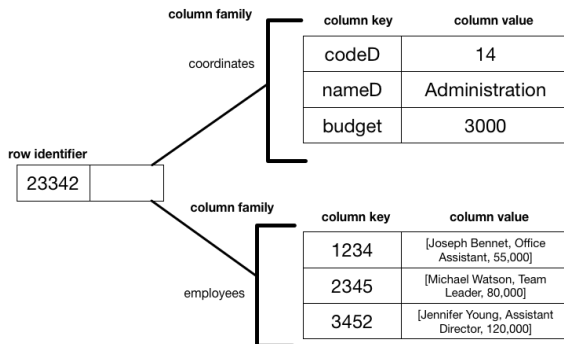
# Column-family Data Model

- Data are modeled as **key-value pairs**.
  - **Key**: row identifier.
  - **Value**: an aggregate, composed of one or more **column families**.
  - **Query**: get a row given its key and the values of its columns.
- **Sharding unit**: a row.
- **Storage unit**: a column family.
- Column-family databases: BigTable, HBase, Cassandra.



# Column-family Data Model

- Definition of a small number of column families.
- As many columns as we need.
- The value of a column can be an aggregate.



# Graph Databases

- DBMS specifically thought to manage and process graphs.
- Two components:
  - **Storage engine**: dictates how the graph is stored.
  - **Processing engine**: dictates how the graph is processed.
- Native storage engine. Storage is tailored to graphs.
- Native processing engine. Operations optimized for graphs.
- Example in details: Neo4j.